

German Traffic Sign Classification

May 19, 2024

```
[1]: pip install opencv-python
```

Requirement already satisfied: opencv-python in c:\apps\anaconda\lib\site-packages (4.9.0.80)

Requirement already satisfied: numpy>=1.21.2 in c:\apps\anaconda\lib\site-packages (from opencv-python) (1.26.4)

Note: you may need to restart the kernel to use updated packages.

```
[2]: !git clone https://bitbucket.org/jadslim/german-traffic-signs
!ls german-traffic-sign
import numpy as np
import matplotlib.pyplot as plt
import keras
from keras.models import Sequential
from keras.optimizers import Adam
from keras.layers import Dense
from keras.layers import Flatten, Dropout
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.layers import Conv2D, MaxPool2D, Dense, Flatten, Dropout
import random
import pickle
import pandas as pd
import cv2
from PIL import Image
```

fatal: destination path 'german-traffic-signs' already exists and is not an empty directory.

'ls' is not recognized as an internal or external command,
operable program or batch file.

```
[3]: from keras.callbacks import LearningRateScheduler, ModelCheckpoint
```

```
%matplotlib inline
np.random.seed(0)
# TODO: Implement load the data here.
with open('german-traffic-signs/train.p', 'rb') as f:
    train_data = pickle.load(f)
with open('german-traffic-signs/valid.p', 'rb') as f:
    val_data = pickle.load(f)
```

```
# TODO: Load test data
with open('german-traffic-signs/test.p', 'rb') as f:
    test_data = pickle.load(f)
```

```
[4]: # Splitting out features and labels
X_train, y_train = train_data['features'], train_data['labels']
X_val, y_val = val_data['features'], val_data['labels']
X_test, y_test = test_data['features'], test_data['labels']
```

```
[5]: #already 4 dimensional
print(X_train.shape)
print(X_val.shape)
print(X_test.shape)
```

```
(34799, 32, 32, 3)
(4410, 32, 32, 3)
(12630, 32, 32, 3)
```

```
[6]: assert(X_train.shape[0] == y_train.shape[0]), "The number of images is not_
↳equal to the number of labels."
assert(X_train.shape[1:] == (32,32,3)), "The dimensions of the images are not_
↳32 x 32 x 3."
assert(X_val.shape[0] == y_val.shape[0]), "The number of images is not equal to_
↳the number of labels."
assert(X_val.shape[1:] == (32,32,3)), "The dimensions of the images are not 32_
↳x 32 x 3."
assert(X_test.shape[0] == y_test.shape[0]), "The number of images is not equal_
↳to the number of labels."
assert(X_test.shape[1:] == (32,32,3)), "The dimensions of the images are not 32_
↳x 32 x 3."
```

```
[7]: data = pd.read_csv('german-traffic-signs/signnames.csv')
print(data)
```

	ClassId	SignName
0	0	Speed limit (20km/h)
1	1	Speed limit (30km/h)
2	2	Speed limit (50km/h)
3	3	Speed limit (60km/h)
4	4	Speed limit (70km/h)
5	5	Speed limit (80km/h)
6	6	End of speed limit (80km/h)
7	7	Speed limit (100km/h)
8	8	Speed limit (120km/h)
9	9	No passing
10	10	No passing for vechiles over 3.5 metric tons
11	11	Right-of-way at the next intersection

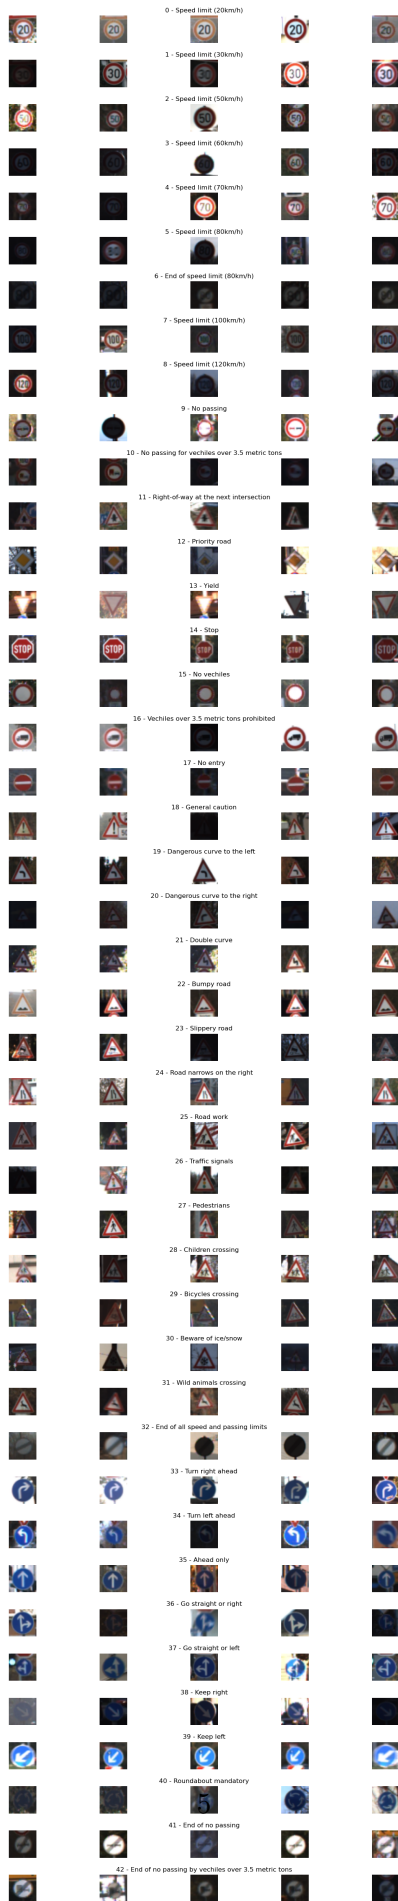
12	12	Priority road
13	13	Yield
14	14	Stop
15	15	No vehicles
16	16	Vehicles over 3.5 metric tons prohibited
17	17	No entry
18	18	General caution
19	19	Dangerous curve to the left
20	20	Dangerous curve to the right
21	21	Double curve
22	22	Bumpy road
23	23	Slippery road
24	24	Road narrows on the right
25	25	Road work
26	26	Traffic signals
27	27	Pedestrians
28	28	Children crossing
29	29	Bicycles crossing
30	30	Beware of ice/snow
31	31	Wild animals crossing
32	32	End of all speed and passing limits
33	33	Turn right ahead
34	34	Turn left ahead
35	35	Ahead only
36	36	Go straight or right
37	37	Go straight or left
38	38	Keep right
39	39	Keep left
40	40	Roundabout mandatory
41	41	End of no passing
42	42	End of no passing by vehicles over 3.5 metric ...

```
[8]: # Assuming 'data' contains the class names
data = pd.read_csv('german-traffic-signs/signnames.csv')
num_of_samples = []
cols = 5
num_classes = 43


# Ensuring that the number of subplots matches the number of classes
fig, axs = plt.subplots(nrows=num_classes, ncols=cols, figsize=(12, 50))
fig.tight_layout()

for i in range(cols):
    for j, row in data.iterrows():
        x_selected = X_train[y_train == j]
        axs[j][i].imshow(x_selected[random.randint(0, (len(x_selected) - 1)), :, 0], cmap=plt.get_cmap('gray'))
```

```
axs[j][i].axis("off")
if i == 2:
    axs[j][i].set_title(str(j) + " - " + row["SignName"])
    num_of_samples.append(len(x_selected))
```



```
[9]: plt.figure(figsize=(12, 4))
plt.bar(range(0, num_classes), num_of_samples)
plt.title("Distribution of the train dataset")
plt.xlabel("Class number")
plt.ylabel("Number of images")
plt.show()
```



```
[10]: import cv2
import pandas as pd


# Defining grayscale and equalization functions
def grayscale(img):
    img_gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    return img_gray

def equalize(img):
    img_equalized = cv2.equalizeHist(img)
    return img_equalized
```

```
[11]: import cv2

plt.imshow(X_train[1000])
plt.axis("off")
print(X_train[1000].shape)
print(y_train[1000])
def grayscale(img):
    img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    return img
```

(32, 32, 3)



```
[12]: import matplotlib.pyplot as plt
import random

def grayscale(img):
    img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    return img

def equalize(img):
    img = cv2.equalizeHist(img)
    return img

def preprocess(img):
    img = grayscale(img)
    img = equalize(img)
    img = img / 255
    return img


X_train = np.array(list(map(preprocess, X_train)))
X_test = np.array(list(map(preprocess, X_test)))
X_val = np.array(list(map(preprocess, X_val)))
```

```

# Display a random grayscale image
plt.imshow(X_train[random.randint(0, len(X_train) - 1)], cmap='gray')
plt.axis('off')
plt.show()
print(X_train.shape)

# Reshape the data
X_train = X_train.reshape(34799, 32, 32, 1)
X_test = X_test.reshape(12630, 32, 32, 1)
X_val = X_val.reshape(4410, 32, 32, 1)

```



```
(34799, 32, 32)
```

```

[13]: print(X_train.shape)
      print(X_test.shape)
      print(X_val.shape)

```

```

(34799, 32, 32, 1)
(12630, 32, 32, 1)
(4410, 32, 32, 1)

```

```

[14]: from tensorflow.keras.utils import to_categorical
      from tensorflow.keras.preprocessing.image import ImageDataGenerator

```



```

# Defining the data augmentation generator
datagen = ImageDataGenerator(
    width_shift_range=0.1,
    height_shift_range=0.1,
    zoom_range=0.2,
    shear_range=0.1,
    rotation_range=10)
# Fitting the data augmentation generator on the training data
datagen.fit(X_train)

```

```

[15]: # Generating augmented data batches
batches = datagen.flow(X_train, y_train, batch_size=20)
X_batch, y_batch = next(batches)

# Visualizing the augmented images
fig, axs = plt.subplots(1, 15, figsize=(20, 5))
fig.tight_layout()

for i in range(15):
    axs[i].imshow(X_batch[i].reshape(32, 32)) # Assuming grayscale images
    axs[i].axis("off")

```



```

[16]: # Converting the labels to one-hot encoded format for train, test, and
      ↪ validation sets
y_train = to_categorical(y_train, 43)
y_test = to_categorical(y_test, 43)
y_val = to_categorical(y_val, 43)

```

```

[19]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
      ↪ Dropout
      from tensorflow.keras.optimizers import Adam

      def modified_model():
          model = Sequential()
          model.add(Conv2D(60, (5, 5), input_shape=(32, 32, 1), activation='relu'))
          model.add(Conv2D(60, (5, 5), activation='relu'))
          model.add(MaxPooling2D(pool_size=(2, 2)))

          model.add(Conv2D(30, (3, 3), activation='relu'))
          model.add(Conv2D(30, (3, 3), activation='relu'))

```

```

model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())
model.add(Dense(500, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(43, activation='softmax'))

model.compile(optimizer=Adam(learning_rate=0.001),
              loss='categorical_crossentropy', metrics=['accuracy'])
return model

model = modified_model()
print(model.summary())

```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_8 (Conv2D)	(None, 28, 28, 60)	1,560
conv2d_9 (Conv2D)	(None, 24, 24, 60)	90,060
max_pooling2d_4 (MaxPooling2D)	(None, 12, 12, 60)	0
conv2d_10 (Conv2D)	(None, 10, 10, 30)	16,230
conv2d_11 (Conv2D)	(None, 8, 8, 30)	8,130
max_pooling2d_5 (MaxPooling2D)	(None, 4, 4, 30)	0
flatten_2 (Flatten)	(None, 480)	0
dense_4 (Dense)	(None, 500)	240,500
dropout_2 (Dropout)	(None, 500)	0
dense_5 (Dense)	(None, 43)	21,543

Total params: 378,023 (1.44 MB)

Trainable params: 378,023 (1.44 MB)

Non-trainable params: 0 (0.00 B)

None

```
[21]: history = model.fit(
    datagen.flow(X_train, y_train, batch_size=50),
    steps_per_epoch=2000,
    epochs=10,
    validation_data=(X_val, y_val),
    shuffle=True
)
```

Epoch 1/10

3/2000 1:20 40ms/step -
accuracy: 0.0322 - loss: 3.7638

C:\Apps\Anaconda\Lib\site-

packages\keras\src\trainers\data_adapters\py_dataset_adapter.py:121:

UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.

self._warn_if_super_not_called()

695/2000 44s 34ms/step -
accuracy: 0.2941 - loss: 2.6157

C:\Apps\Anaconda\Lib\contextlib.py:158: UserWarning: Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches. You may need to use the `.repeat()` function when building your dataset.

self.gen.throw(typ, value, traceback)

2000/2000 27s 13ms/step -
accuracy: 0.4288 - loss: 2.0595 - val_accuracy: 0.9084 - val_loss: 0.2926

Epoch 2/10

2000/2000 30s 15ms/step -
accuracy: 0.8171 - loss: 0.5888 - val_accuracy: 0.9671 - val_loss: 0.1168

Epoch 3/10

2000/2000 31s 16ms/step -
accuracy: 0.8921 - loss: 0.3463 - val_accuracy: 0.9737 - val_loss: 0.0899

Epoch 4/10

2000/2000 34s 17ms/step -
accuracy: 0.9219 - loss: 0.2474 - val_accuracy: 0.9751 - val_loss: 0.0974

Epoch 5/10

2000/2000 34s 17ms/step -
accuracy: 0.9381 - loss: 0.2014 - val_accuracy: 0.9814 - val_loss: 0.0661

Epoch 6/10

2000/2000 34s 17ms/step -
accuracy: 0.9448 - loss: 0.1736 - val_accuracy: 0.9866 - val_loss: 0.0512


Epoch 7/10

2000/2000 34s 17ms/step -

```
accuracy: 0.9498 - loss: 0.1611 - val_accuracy: 0.9848 - val_loss: 0.0532
Epoch 8/10
2000/2000          36s 18ms/step -
accuracy: 0.9594 - loss: 0.1301 - val_accuracy: 0.9807 - val_loss: 0.0753
Epoch 9/10
2000/2000          39s 19ms/step -
accuracy: 0.9606 - loss: 0.1258 - val_accuracy: 0.9887 - val_loss: 0.0335
Epoch 10/10
2000/2000          35s 18ms/step -
accuracy: 0.9629 - loss: 0.1217 - val_accuracy: 0.9907 - val_loss: 0.0356
```


```
[46]: import matplotlib.pyplot as plt

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.legend(['Training', 'Validation']) # Corrected legend syntax
plt.title('Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss') # Added ylabel for clarity
plt.show()
```



```
[47]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.legend(['training', 'validation'])
plt.title('Accuracy')
plt.xlabel('epoch')
```

```
[47]: Text(0.5, 0, 'epoch')
```



```
[48]: # Evaluating model on test data
score = model.evaluate(X_test, y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
Test score: 0.10784535109996796
Test accuracy: 0.9752969145774841
```

```
[49]: # Evaluating model on training data
train_score = model.evaluate(X_train, y_train, verbose=0)
print('Training score:')
```

```

print(' Loss:', train_score[0])
print(' Accuracy:', train_score[1])

# Evaluating model on validation data
val_score = model.evaluate(X_val, y_val, verbose=0)
print('Validation score:')
print(' Loss:', val_score[0])
print(' Accuracy:', val_score[1])

```

Training score:

Loss: 0.016071079298853874

Accuracy: 0.9951435327529907

Validation score:

Loss: 0.035643890500068665

Accuracy: 0.9907029271125793

```

[42]: # Predicting internet image
import requests
from PIL import Image
import numpy as np
import cv2

# URL of the image
url = 'https://c8.alamy.com/comp/J2MRAJ/
↳german-road-sign-bicycles-crossing-J2MRAJ.jpg'

# Downloading and loading the image
r = requests.get(url, stream=True)
img = Image.open(r.raw)
plt.imshow(img, cmap=plt.get_cmap('gray'))

# Converting the image to grayscale, resize, and preprocess
img = np.asarray(img)
img = cv2.resize(img, (32, 32))
img = preprocess(img)
img = img.reshape(1, 32, 32, 1)

# Predicting the class probabilities
predictions = model.predict(img)

# Getting the index with the highest probability
predicted_class = np.argmax(predictions)

print("Predicted class:", predicted_class)

```

1/1 0s 18ms/step

Predicted class: 29



THE END. THANK YOU!

[]:

[]:

[]:

[]: